

Amazon DocumentDB

Amazon **DocumentDB** is a fully managed, scalable NoSQL database service designed to store and query JSON-like document data. It's built to be compatible with MongoDB APIs, allowing developers to migrate or build document-based applications easily without managing infrastructure.

Key Features:

- **MongoDB Compatibility:** Supports MongoDB workloads and drivers with minimal code changes.
- **Document-Oriented Storage:** Stores semi-structured data in flexible JSON-like format.
- **Separation of Compute and Storage:** Enables independent scaling and high availability.
- **Automatic Backups and Point-in-Time Recovery:** Built-in data protection without manual setup.
- **High Availability:** Supports replication across multiple AZs with automatic failover.
- **Fully Managed:** No server provisioning, patching, or maintenance needed.
- **Security Integration:** Works with VPC, KMS encryption, and IAM for fine-grained access control.

Use Case:

Ideal for content management systems, product catalogs, and user profile stores where data structure is flexible and may vary between records — especially useful when teams want MongoDB-like functionality in a managed AWS-native solution.

Amazon Neptune + Neptune Streams

Amazon Neptune is a fully managed graph database service optimized for storing and querying highly connected data using graph models like property graph (Gremlin) and RDF (SPARQL). It allows you to efficiently navigate relationships between data elements.

Key Features:

- Supports Gremlin and SPARQL: Use industry-standard graph query languages for both property graph and RDF.
- Highly Available and Durable: Replication across AZs with automatic failover.
- High Performance: Optimized for low-latency graph traversal queries.
- Fully Managed: No need to provision, patch, or manage database infrastructure.
- Encryption at Rest and in Transit: Integrated with AWS KMS.
- IAM Access Control: Fine-grained access to Neptune APIs.
- Integration with Neptune Streams: Streams capture changes in your graph (add/delete nodes or edges) to enable event-driven architectures.

Neptune Streams:

Allows you to capture change logs (add/remove of vertices/edges) in real-time and process them with AWS services like Lambda or Kinesis — ideal for analytics, auditing, or triggering business logic.

Use Case:

Perfect for use cases like fraud detection, recommendation engines, social networking graphs, and knowledge graphs — where relationships between data are as important as the data itself. Neptune Streams is especially useful when you need to react in real time to graph changes (e.g., notifying users of new connections).

Keyspaces (for Apache Cassandra)

Amazon **Keyspaces** is a fully managed serverless database service that supports Apache Cassandra workloads. It enables you to run Cassandra workloads on AWS without managing infrastructure, scaling seamlessly as traffic grows.

Key Features:

- **Cassandra-Compatible:** Use existing Cassandra drivers and query language (CQL) with no changes.
- **Serverless and Scalable:** Automatically scales based on application traffic — no capacity planning.
- **High Availability:** Replicated across multiple AZs for resilience.
- **Secure:** Encryption at rest using KMS and VPC integration for private access.
- **Pay-per-Request Pricing:** Only pay for what you use, ideal for variable workloads.
- **Integration with IAM:** Control access with fine-grained IAM policies.
- **Monitoring and Logging:** Integrated with CloudWatch for observability.

Use Case:

Ideal for developers already using Cassandra who want to migrate to a managed, scalable solution — for example, storing sensor data in an IoT application, managing user profiles, or handling session data for mobile apps.

Amazon QLDB (Quantum Ledger Database)

Amazon **QLDB** is a fully managed ledger database designed to provide a transparent, immutable, and cryptographically verifiable transaction log. It tracks every application data change and maintains a complete and verifiable history over time.

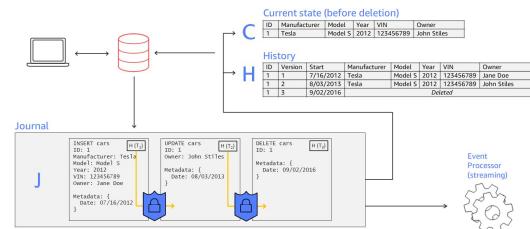
Key Features:

- **Immutable Ledger:** Once data is entered, it cannot be changed or deleted.
- **Cryptographic Verification:** Built-in hash chains allow cryptographic validation of changes.
- **SQL-like Query Language (PartiQL):** Enables easy querying without needing to learn blockchain-specific syntax.
- **Managed Service:** No need to manage hardware, provisioning, or maintenance.
- **High Performance:** Optimized for high throughput and low-latency workloads.
- **Serverless:** Automatically scales up/down and handles infrastructure.

Use Case:

Useful for systems that need a trusted audit trail — such as financial transactions, supply chain tracking, or regulatory compliance logs — without the complexity of blockchain.

Amazon QLDB: the journal is the database



Amazon Timestream

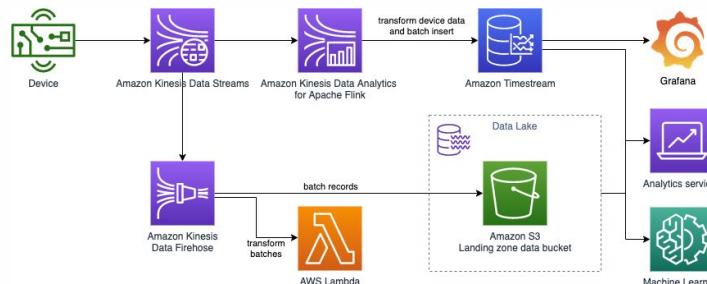
Amazon **Timestream** is a serverless, time series database designed for collecting, storing, and analyzing time-stamped data at scale. It's optimized for use cases like IoT sensor data, operational metrics, and application monitoring.

Key Features:

- **Time Series Optimized:** Automatically organizes data by time, making queries faster and storage more efficient.
- **Serverless:** No server provisioning; it automatically scales with demand.
- **Data Lifecycle Management:** Automatically moves old data to a cost-optimized storage tier.
- **SQL Support:** Use familiar SQL syntax to analyze time series data.
- **High Availability:** Built-in replication and backup.
- **Integration:** Works seamlessly with AWS services like IoT Core, Lambda, Kinesis, and Grafana.

Use Case:

Ideal for tracking metrics over time, such as CPU utilization of EC2 instances, temperature readings from IoT devices, or real-time analytics dashboards for applications.



DB Exam Questions

A financial services company needs to implement a ledger-like database for tracking all transactions with cryptographic verification and immutability. Additionally, they are building a graph-based recommendation engine that requires fast traversals of complex relationships.

Which two AWS database services should they use?

- A. Amazon DocumentDB
- B. Amazon Neptune
- C. Amazon QLDB
- D. Amazon Keyspaces
- E. Amazon Timestream

A startup is building a highly available SaaS application that requires a relational database, automatic replication across multiple AZs, failover within seconds, and compatibility with MySQL and PostgreSQL. Which AWS database service best meets these requirements?

- A. Amazon Timestream
- B. Amazon DocumentDB
- C. Amazon Keyspaces
- D. Amazon Aurora

Amazon Athena

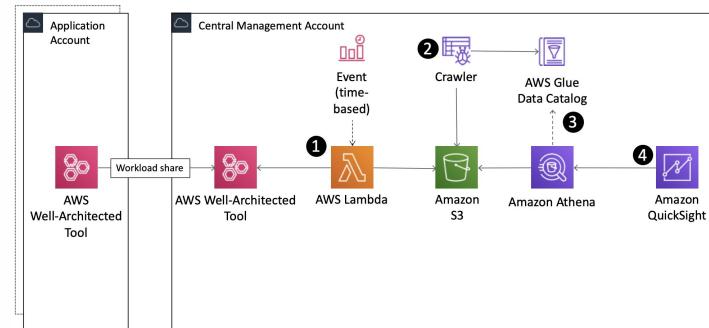
Amazon **Athena** is a serverless, interactive query service that allows you to use SQL to analyze data directly in S3 without setting up any infrastructure.

Key Features:

- Supports standard SQL
- Serverless - No need to manage compute resources
- Integrated with AWS Glue Data Catalog
- Charges per query and scanned data
- Uses Presto under the hood for query execution

Use Cases:

Ad-hoc querying on S3 datasets, quick analysis on logs, reports, clickstreams, or CSV/Parquet files.



Amazon Athena – Performance Improvement

Athena performance can be significantly improved by optimizing how data is stored and queried:

1. Partitioning – Split data based on logical keys (e.g., date, region) so that Athena only scans the relevant partitions instead of the entire dataset.
2. Compression – Use compressed formats like GZIP or Snappy to reduce the amount of data read and lower query costs.
3. Columnar Storage – Store data in columnar formats like Parquet or ORC to allow Athena to scan only needed columns, improving speed and efficiency.
4. Optimized File Sizes – Avoid too many small files or overly large files. Aim for optimal file sizes (typically between 128MB to 1GB) to ensure efficient processing and cost balance.

Use Case:

Athena is ideal for querying large datasets in S3 when combined with these performance optimizations, reducing both query time and cost.

Amazon Athena – Federated Query

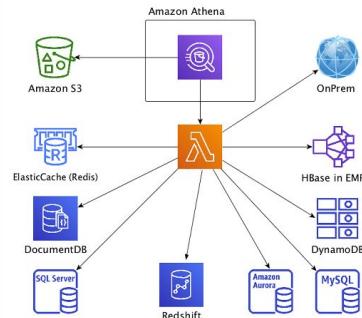
Athena Federated Query allows you to run SQL queries across multiple data sources—not just Amazon S3. It uses data source connectors (powered by AWS Lambda) to connect to external sources like:

- Relational databases (e.g., RDS, Aurora)
- NoSQL databases (e.g., DynamoDB)
- Third-party sources (e.g., Snowflake, MySQL on-prem)

With federated queries, you can join and query data from different locations using standard SQL in Athena, without needing to move or copy the data.

Use Case:

Enables data analysts to combine real-time data from operational databases with historical data stored in S3, all within one query.



Amazon Redshift – Overview & Cluster

Amazon Redshift is AWS's fully managed data warehouse solution optimized for OLAP (Online Analytical Processing) and large-scale analytics workloads. It's designed for fast querying and analysis over petabytes of structured data.

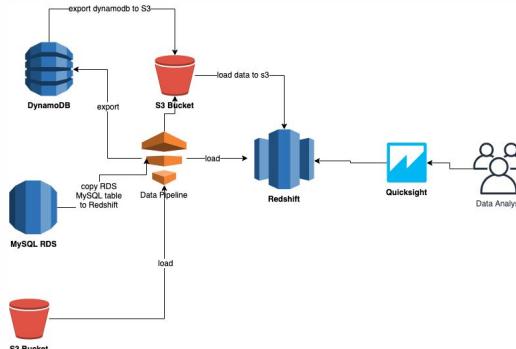
Redshift organizes its environment into clusters, which are collections of nodes. Each cluster has:

1. **Leader Node:** Manages query coordination and planning.
2. **Compute Nodes:** Store data and execute queries.

You can start with a single-node cluster or a multi-node cluster for high performance and scalability. Redshift also supports RA3 nodes with managed storage and AQUA (Advanced Query Accelerator) to offload some operations for faster performance.

Use Case:

Run complex analytics and BI (business intelligence) dashboards that query data across massive volumes with sub-second response times.



Amazon Redshift – Snapshots & Disaster Recovery

Amazon Redshift provides automated and manual snapshots to backup your data:

Automated Snapshots:

Enabled by default. Created every 8 hours or after 5 GB of changes. Retained for a default of 1 day (can be customized).

Manual Snapshots:

User-initiated and kept until explicitly deleted.

All snapshots are stored in Amazon S3, and you can restore a snapshot to a new Redshift cluster in any region, which makes it ideal for Disaster Recovery strategies.

You can also share manual snapshots across AWS accounts and regions.

Use Case:

Implementing cross-region DR by creating manual snapshots and restoring them in a different region in case of an outage or failure.

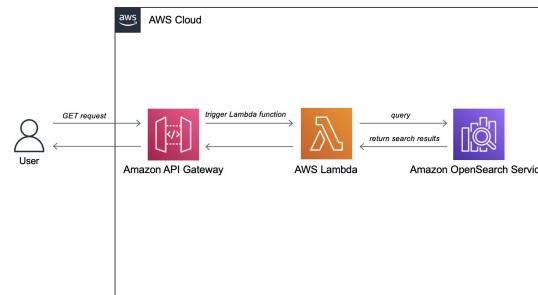
Amazon OpenSearch Service

Amazon OpenSearch is a fully managed service that makes it easy to deploy, operate, and scale OpenSearch clusters in the AWS Cloud.

- Search & Analytics Engine: Used for real-time log analytics, full-text search, application monitoring, and clickstream analytics.
- Ingesting Data: You can ingest data from Amazon Kinesis Data Firehose, Logstash, AWS Lambda, or directly via the OpenSearch APIs.
- Visualizing Data: Integrated with OpenSearch Dashboards (formerly Kibana) for powerful data visualizations.
- Security: Supports encryption at rest and in transit, VPC access, IAM-based access control, and fine-grained user permissions.

Use Case:

Real-time monitoring of application logs and search functionality in enterprise applications.



Amazon EMR (Elastic MapReduce)

Amazon **EMR** is a managed big data platform that allows you to process vast amounts of data using open-source frameworks such as Apache Hadoop, Spark, Hive, Presto, and more.

Node Types:

- **Master Node**: Manages the cluster, tracks job status, and coordinates task distribution.
- **Core Node**: Runs Hadoop tasks and stores data using HDFS (replicated).
- **Task Node**: Only runs tasks (no HDFS storage), used for parallel processing.

Key Points:

- You can customize clusters using bootstrap actions or EMR steps.
- Supports transient (on-demand) and long-running clusters.
- Can use Spot Instances for cost savings.
- Integrates with S3, DynamoDB, Redshift, RDS, and more.
- Auto-scaling and auto-healing supported.
- You can install additional applications using the EMR console or CLI.

Use Case:

Running distributed big data processing tasks like log analysis, ETL pipelines, or machine learning at scale.

Amazon QuickSight

Amazon **QuickSight** is a fully managed BI (Business Intelligence) service that enables you to create interactive dashboards and reports from multiple data sources. It uses the SPICE engine (Super-fast, Parallel, In-memory Calculation Engine) to accelerate data processing.

Key Features:

- SPICE engine enables fast, in-memory data analysis without hitting the data source each time.
- Integrates with S3, Athena, Redshift, RDS, Snowflake, MySQL, and more.
- IAM-based access control for secure data sharing.
- Dashboards can be embedded in apps and websites.
- Supports natural language querying via the "Q" feature.
- Flexible pricing: Pay-per-session or monthly licensing.

Use Case:

Ideal for organizations that want real-time insights across different data sources and need to provide visual dashboards for executives, analysts, or end-users.



AWS Glue

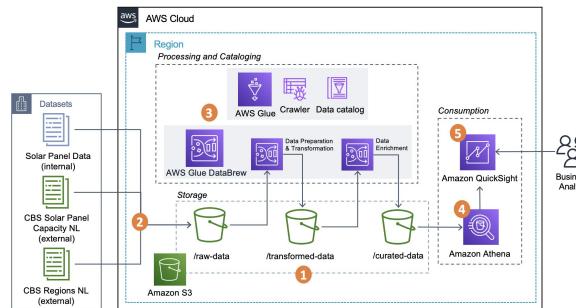
AWS **Glue** is a serverless data integration service used to prepare, transform, and load (ETL) data for analytics, machine learning, and application development. It automates much of the effort in building, maintaining, and running ETL jobs.

Key Features:

- Serverless ETL: No infrastructure to manage.
- Crawler: Automatically discovers data schema and adds it to the Glue Data Catalog.
- Supports Python and Scala for custom ETL scripts.
- Glue Studio: Visual interface for building ETL pipelines.
- Glue DataBrew: No-code tool to clean and normalize data.
- Job scheduling and dependency management included.

Use Case:

Great for organizations that need to transform raw data from various sources (S3, RDS, JDBC, etc.) into structured formats ready for querying and analytics.



AWS Lake Formation

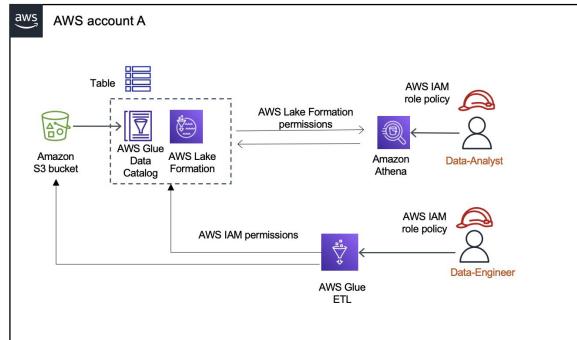
AWS **Lake Formation** is a service that simplifies building, securing, and managing data lakes on Amazon S3. It allows you to centralize, organize, and secure your data for analytics and ML.

Key Features:

- Data lake creation: Automates importing data from S3, RDS, or other sources.
- Fine-grained access control: Centralizes data permissions across services like Athena, Redshift, and Glue.
- Data catalog integration: Uses AWS Glue Data Catalog to manage metadata.
- Data transformation: Integrates with Glue for preparing and transforming data.
- Row-level security & column masking for enhanced data governance.

Use Case:

Ideal for enterprises building a centralized, secure data repository where multiple services and users can query the same data securely and efficiently.



Kinesis Data Analytics (SQL + Apache Flink)

Kinesis Data Analytics is a fully managed service for analyzing streaming data in real time using either SQL or Apache Flink.

Modes of Operation:

1. **SQL Mode:**

- Simple SQL queries to filter, aggregate, or join real-time data.
- Ideal for users who don't want to manage infrastructure or write Java/Scala code.

2. **Apache Flink Mode:**

- For advanced stream processing using the Apache Flink framework.
- Supports windowed operations, stateful processing, and complex event patterns.

Supported Sources and Sinks:

- Works with Kinesis Data Streams, Kinesis Data Firehose, and Apache Kafka (via MSK).
- Can output to S3, Redshift, OpenSearch, or another Kinesis stream.

Use Case:

Perfect for real-time dashboards, alerting systems, and live analytics on streaming data without needing to manage stream processing infrastructure.

Amazon MSK (Managed Streaming for Apache Kafka)

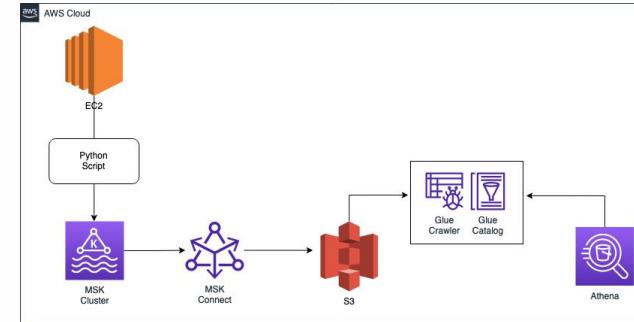
Amazon **MSK** is a fully managed service that makes it easy to build and run applications using Apache Kafka for streaming data. It takes away the complexity of setting up, securing, and maintaining Kafka clusters.

Key Features:

- Fully managed: No need to manually install, patch, or operate Kafka or Zookeeper.
- Secure: Integrated with AWS IAM, VPC, encryption in transit and at rest.
- Highly available: MSK handles replication and automatic recovery across multiple AZs.
- Scalable: Can scale up by adding brokers and partitions.
- Monitoring: Built-in integration with CloudWatch for metrics and logs.

Use Case:

Used for real-time data pipelines, event-driven architectures, or log processing systems — for example, collecting and processing logs or clickstream data from thousands of websites or IoT devices.



Data & Analytics Exam Questions

A company needs a fully managed search and log analytics service that can index structured and unstructured logs, perform full-text search, and generate dashboards for operational monitoring.

Which AWS service should they use?

- A. Amazon OpenSearch Service
- B. Amazon Redshift
- C. AWS Glue
- D. Amazon MSK

A data analyst wants to run ad-hoc SQL queries directly on raw data stored in Amazon S3 without setting up or managing any infrastructure. The solution should be cost-effective and serverless, charging only for the queries run.

Which service is the most suitable?

- A. Amazon EMR
- B. Amazon Glue
- C. Amazon Athena
- D. Amazon Redshift

AWS Machine Learning

1. **Rekognition** - Detects faces, labels objects, recognizes celebrities.

Keywords: image analysis, facial detection

2. **Transcribe** - Converts audio to text, useful for subtitles or transcripts.

Keywords: speech-to-text, audio transcription

3. **Polly** - Converts text into lifelike speech.

Keywords: text-to-speech, voice synthesis

4. **Translate** - Provides real-time language translation.

Keywords: language translation, real-time translation

5. **Lex** - Builds conversational chatbots with voice or text.

Keywords: chatbot, natural language

6. **Connect** - Cloud-based contact center service.

Keywords: call center, customer service

7. **Comprehend** - Performs NLP to extract insights from text.

Keywords: NLP, sentiment analysis

8. **SageMaker** - Fully managed ML service to build, train, and deploy models.

Keywords: train model, machine learning

9. **Forecast** - Uses ML to deliver time-series forecasting.

Keywords: predict demand, forecasting

10. **Kendra** - ML-powered search engine for enterprise data.

Keywords: intelligent search, enterprise search

11. **Personalize** - Provides real-time recommendations (e.g., movies, products).

Keywords: user recommendations, personalization

12. **Textract** - Extracts text, tables, and forms from scanned documents.

Keywords: OCR, document analysis

ML Exam Questions

A retail company is implementing a smart surveillance system to detect unusual activity in its stores. As part of the solution, they want to identify objects and detect people's faces from security camera footage in real time. Which AWS service provides built-in support for image and video analysis, including facial detection and object labeling?

- A. Polly
- B. Rekognition
- C. Comprehend
- D. Kinesis Data Analytics

A healthcare startup wants to convert doctor-patient conversation recordings into text for medical records. The solution should be able to handle audio input and return accurate, timestamped transcripts. Which AWS service is the most suitable for this requirement?

- A. Translate
- B. Textract
- C. Transcribe
- D. Lex

Amazon CloudWatch

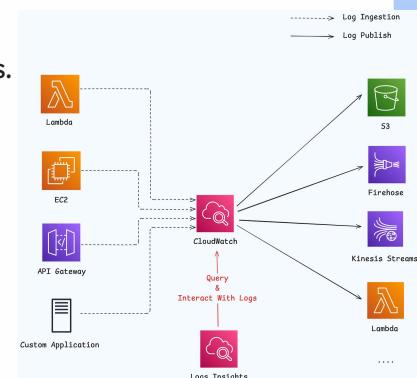
Amazon **CloudWatch** is a powerful monitoring and observability service designed to help DevOps engineers, system administrators, and developers monitor AWS resources and the applications they run in real time. It collects and tracks metrics, collects and monitors log files, and sets alarms to automatically react to changes in your AWS environment. CloudWatch can monitor AWS services like EC2, RDS, Lambda, and custom metrics from on-premises servers or applications.

Key Features:

- **Metrics Collection:** Tracks performance indicators (like CPU usage, disk reads/writes, network traffic).
- **Alarms:** You can set thresholds and be alerted via SNS or trigger automated actions.
- **Dashboards:** Visualize data from multiple sources in a single place.
- **Log Collection:** Native integration with CloudWatch Logs for real-time log collection and analysis.
- **Event-driven Automation:** Integrated with EventBridge to trigger actions based on specific events.

Common Use Cases:

- Detect underperforming EC2 instances by analyzing CPU utilization.
- Alert when S3 bucket size exceeds a specific threshold.
- Automate recovery or scaling based on metrics (e.g., autoscale EC2 when average CPU > 70%).



CloudWatch Logs Insights

CloudWatch Logs Insights is a powerful log analytics tool that allows you to interactively search and analyze log data in real time using a purpose-built query language. It helps you troubleshoot application and infrastructure issues faster by running queries on logs stored in CloudWatch Logs without the need to move them to external systems.

Key Features:

- **Query Language:** Similar to SQL, lets you filter, parse, and aggregate logs efficiently.
- **Real-time Insights:** Run queries on recent log data or over custom time ranges.
- **Visualization:** View query results in tables and graphs.
- **Performance:** Scales automatically with the amount of log data; optimized for high-speed execution.
- **Integrated with Dashboards:** Query results can be pinned directly to CloudWatch dashboards.

Common Use Cases:

- Analyze Lambda execution logs to find errors and durations.
- Troubleshoot ECS/EKS application issues by searching container logs.
- Monitor login attempts, application exceptions, or failed requests across services.

CloudWatch Logs Aggregation – Multi-Account & Multi-Region + Subscriptions

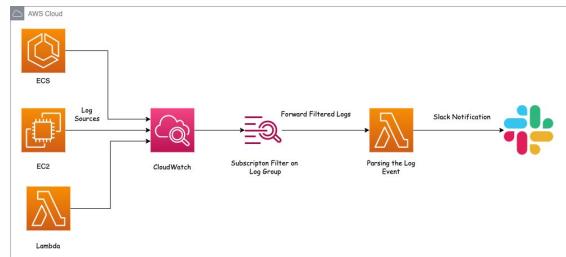
CloudWatch Logs **Aggregation** allows centralizing log data from multiple AWS accounts and regions into a single destination for centralized visibility and monitoring. This is essential for organizations operating in multi-account or multi-region environments.

Key Features:

- **Cross-account and cross-region log centralization:** Use subscriptions to stream logs from different accounts or regions into a central account.
- **Log Subscriptions:** Define a subscription filter on a log group to send matching logs in real-time to a destination—like a Kinesis stream, Lambda, or a central CloudWatch log group.
- **Log Destination Resource:** Acts as a target for log subscriptions, typically used in central logging architectures.
- **IAM Permissions:** Destination accounts must grant permissions to allow log delivery from source accounts.

Benefits:

1. Unified monitoring across all workloads, regardless of account or region.
2. Easier compliance and audit trail management.
3. Streamlined security operations (e.g., detect anomalies across all environments from a single place).



CloudWatch Logs Agent & Unified CloudWatch Agent

To collect logs and system-level metrics from EC2 instances (and on-premises servers), AWS provides two main agents: the older CloudWatch Logs Agent and the more modern Unified CloudWatch Agent.

CloudWatch Logs Agent:

Designed only to push logs (e.g.,
/var/log/messages, custom app
logs) to CloudWatch Logs.

Configured using a config.json file.

Simpler but limited—does not
support metrics collection.

Older tool, mostly used in legacy
systems.

Unified CloudWatch Agent:

A newer, single agent that supports both logs and metrics collection.

Allows collecting:
OS-level metrics (CPU, memory, disk, network).
Custom metrics.
Log files.

Configured using amazon-cloudwatch-agent.json or via the SSM parameter store.
Works across platforms: Linux, Windows, on-prem, and EC2.
Supports centralized deployment using Systems Manager (SSM).
Recommended for new deployments.

CloudWatch Alarms + Targets + Composite Alarms

CloudWatch Alarms are used to monitor metrics and respond to changes in your AWS resources. They help you automate responses and alerting based on thresholds you define.

Standard CloudWatch Alarms:

Monitor a single metric (e.g., CPUUtilization > 80% for 5 minutes).

Trigger actions when the metric is AboveThreshold, BelowThreshold, or Missing.

Actions include:

- Sending an SNS notification (email, SMS, Lambda, etc.).
- Auto-scaling actions (scale out/in EC2).
- EC2 recovery or reboot.

Alarm States:

1. **OK** – Metric within the expected range.
2. **ALARM** – Metric crossed the defined threshold.
3. **INSUFFICIENT DATA** – Not enough data (e.g., newly created instance).

Composite Alarms:

Evaluate multiple alarms together using logical AND/OR operators.

Do not incur additional charges for evaluation.

Targets/Actions:

Common alarm actions:

- Triggering SNS Topics to notify operations teams.
- Invoking Lambda functions to take corrective actions.
- EC2 actions: recover, stop, terminate, or reboot.
- Auto Scaling: increase or decrease instance count.

Use cases:

Reduce noise by only alerting if multiple conditions are met.

Example: Trigger alert only if both CPU > 80% AND memory usage > 90%.

EC2 Instance Recovery

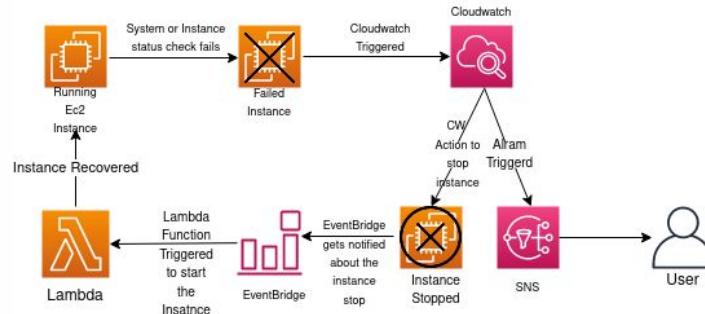
EC2 Instance Recovery is an automated mechanism provided by CloudWatch to recover an impaired EC2 instance due to hardware or system issues. It stops and restarts the instance on healthy hardware within the same Availability Zone, retaining the same instance ID, Elastic IP, and data stored on EBS.

Key Points:

- Initiated via CloudWatch alarm on StatusCheckFailed_System.
- Preserves instance metadata and attached volumes.
- Suitable for production workloads that require minimal downtime.
- Does not fix software/application-level issues.

Use Case:

- A critical web server running on EC2 crashes due to hardware failure. CloudWatch triggers a recovery action, restarting the instance automatically on healthy hardware—minimizing service disruption and avoiding manual intervention.



EventBridge

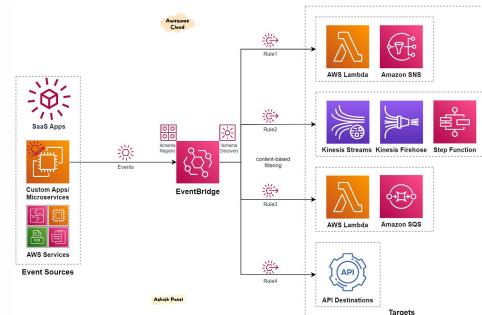
Amazon EventBridge is a serverless event bus service that connects application data from your own applications, AWS services, and third-party SaaS applications. It allows you to build event-driven architectures by routing events based on rules to AWS targets such as Lambda, SNS, SQS, Step Functions, and more.

Key Points:

- Ingests events from over 200 AWS services and SaaS providers.
- Uses rules to match incoming events and route them to specific targets.
- Supports event filtering and transformation before delivery.
- Natively integrates with schema discovery and schema registry for easier development.
- Suitable for decoupling microservices and building loosely coupled applications.

Use Case:

When a new file is uploaded to an S3 bucket, EventBridge routes this event to a Lambda function that processes the file and stores metadata in DynamoDB—without requiring polling or manual triggers.



CloudWatch Insights and Operational Visibility

CloudWatch **Insights** is a suite of tools within Amazon CloudWatch designed to give deep operational visibility across AWS environments. It includes various modules:

1. **CloudWatch Container Insights**: Provides detailed metrics and logs for containerized applications running on ECS, EKS, and Kubernetes on EC2 or Fargate. Requires a CloudWatch agent for Kubernetes clusters.
2. **CloudWatch Lambda Insights**: Offers detailed performance metrics and diagnostics for Lambda functions, including CPU time, memory usage, and cold start data.
3. **CloudWatch Contributor Insights**: Helps identify the "Top-N" contributors to system anomalies or high-usage patterns based on log data.
4. **CloudWatch Application Insights**: Automatically sets up dashboards and monitoring for common application stacks and services to detect issues like memory leaks or latency spikes.

These tools provide developers, DevOps, and system operators with end-to-end observability and help in pinpointing performance bottlenecks and improving system reliability.

Use Case:

A developer troubleshoots high latency in a Lambda-based API by using Lambda Insights to discover frequent cold starts and then uses Application Insights to correlate the issue with high memory usage in a connected RDS instance.

CloudTrail + Insights + Event Retention

AWS **CloudTrail** is a service that enables governance, compliance, and operational and risk auditing by recording API activity across your AWS account. It captures events from the AWS Management Console, AWS SDKs, command-line tools, and other AWS services.

CloudTrail Insights: An advanced feature that automatically detects unusual API activity patterns in your account. This helps identify security threats, misconfigurations, or operational issues.

Event Retention: CloudTrail logs are stored in S3, and you can define retention policies. You can also integrate logs with CloudWatch Logs for real-time monitoring.

CloudTrail is region-aware, and you can create multi-region trails or organization-wide trails with AWS Organizations for centralized auditing.

Use Case:

A security team notices increased StartInstances API calls during non-business hours. CloudTrail Insights flags this behavior, helping the team quickly identify a compromised IAM user and take corrective action.

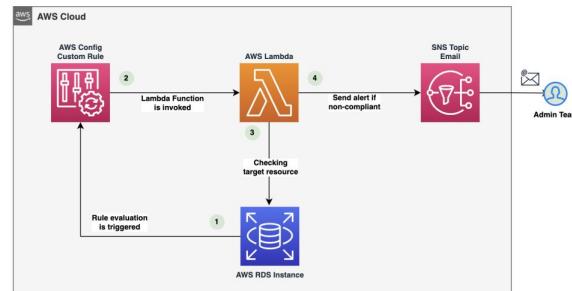
AWS Config

AWS **Config** is a service that enables you to assess, audit, and evaluate the configurations of your AWS resources. It continuously monitors and records AWS resource configurations and allows you to automate compliance checks against desired configurations.

- **Configuration History:** AWS Config maintains a history of all configuration changes over time, enabling you to trace how a resource was changed and by whom.
- **Rules and Remediation:** You can create AWS Config rules to evaluate resource compliance (e.g., EC2 instances must be in a specific VPC) and automatically trigger remediation actions using Systems Manager Automation.
- **Integration with AWS Organizations:** Config can be set up at the organization level for centralized compliance and auditing across multiple accounts.

Use Case:

A compliance team wants to ensure that all S3 buckets are encrypted. AWS Config continuously evaluates this rule and automatically flags any bucket that is not compliant, optionally triggering an automation to enable encryption.



CloudWatch vs CloudTrail vs AWS Config

Amazon CloudWatch

Focus: Performance monitoring

Collects metrics, logs, and events from AWS resources and custom applications.

Used for real-time monitoring, alerting, and dashboards.

Example: Detecting high CPU utilization on an EC2 instance.

AWS CloudTrail

Focus: API activity logging

Records who did what in your AWS account, including console actions and API calls.

Crucial for audit trails, security analysis, and compliance.

Example: Tracking who terminated an EC2 instance or modified a security group.

AWS Config

Focus: Configuration compliance

Continuously records resource configurations and evaluates them against compliance rules.

Enables historical tracking of how a resource changed over time.

Example: Checking if all EBS volumes are encrypted and reporting non-compliant ones.

1. Use CloudTrail to detect unauthorized access or suspicious activity,
2. Use CloudWatch to monitor system health and performance,
3. Use AWS Config to ensure resources follow security policies like encryption and tagging standards.

Monitoring Exam Questions

A company wants to ensure that all Amazon EC2 instances launched in its environment meet compliance standards. The requirements include using a customer-managed KMS key for encryption and tagging each instance with an "Environment" key. Additionally, the security team must be alerted automatically if an instance becomes non-compliant. Which solution meets these requirements with minimal operational overhead?

- A. Integrate Amazon CloudWatch with AWS Lambda to monitor EC2 tags and encryption status and send alerts.
- B. Use AWS CloudTrail to track API activity and manually review logs for non-compliant instances.
- C. Use Amazon Inspector to scan EC2 instances for tag and encryption compliance and report issues.
- D. Create custom AWS Config rules to evaluate encryption and tags, and use Amazon SNS to notify the security team of non-compliant resources.

An operations team needs to investigate why a production EC2 instance was unexpectedly stopped yesterday afternoon. The team must identify which IAM user or role initiated the stop action and when it occurred. Which solution will provide this information in the most efficient and reliable way?

- A. Query AWS CloudTrail logs to find the StopInstances API call and identify the IAM user or role that initiated it.
- B. Use Amazon CloudWatch to view recent EC2 metrics and determine when CPU utilization dropped.
- C. Query AWS CloudTrail logs to find the StopInstances API call and identify the IAM user or role that initiated it.
- D. Use AWS Trusted Advisor to check for underutilized instances and stop history.

AWS Organizations

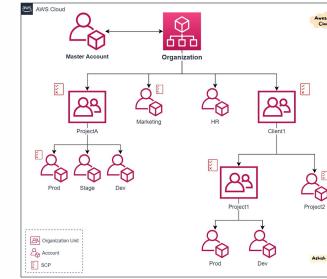
AWS **Organizations** is a service that allows you to centrally manage multiple AWS accounts. With it, you can automate account creation, group accounts into organizational units (OUs), and apply governance using Service Control Policies (SCPs). Organizations support consolidated billing, making it easier to track and optimize costs across your environment. It is especially useful in enterprise setups where workloads are separated by departments, teams, or environments (e.g., dev, test, prod).

Key Capabilities:

- Centralized billing and cost control
- Policy-based governance with SCPs
- Organizational Units for account grouping
- Programmatic account creation and invitation
- Integration with AWS IAM Identity Center and Control Tower

Use Case:

A company wants to separate billing and access for different departments (Finance, Marketing, DevOps). They use AWS Organizations to create one account per department, manage them from a master account, and apply policies to restrict service access per department.



Service Control Policies (SCP)

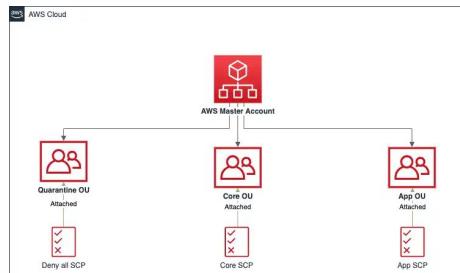
Service Control Policies (SCPs) are a core feature of AWS Organizations that allow you to define the maximum permissions for accounts or organizational units (OUs) in your organization. SCPs don't grant permissions by themselves—they act as a guardrail, limiting what IAM users and roles can do, even if an IAM policy allows it. SCPs apply to all users and roles in an account, including the root user.

Key Concepts:

- SCPs define permission boundaries at the account or OU level.
- SCPs can be deny-only or allow-specific actions.
- The default policy is FullAWSAccess, which allows everything unless restricted.
- SCPs do not override resource-based or identity-based policies, but limit their effective scope.

Use Case:

An organization wants to ensure that no user, in any account under the "Dev" OU, can delete S3 buckets—even if their IAM policy allows it. An SCP is attached to the "Dev" OU with an explicit Deny for s3:DeleteBucket, ensuring consistent control.



IAM Conditions + Permission Boundaries

IAM Conditions and Permission Boundaries are advanced mechanisms to fine-tune access control in AWS:

- **IAM Conditions:** These are key-value pairs defined inside IAM policies that allow you to enforce context-aware rules—like allowing access only from a specific IP address, or during certain hours, or requiring MFA.
- **Permission Boundaries:** These are special policies that define the maximum permissions a user or role can be granted, even if their identity-based policy allows more. They're like a “cap” on what IAM entities can do.

Key Differences:

1. Conditions filter when and how permissions apply.
2. Boundaries restrict what permissions are available at most.

Use Case:

An organization wants a junior developer to only manage EC2 instances within a specific region, and only if they are tagged with Environment=Dev. The IAM policy has conditions based on region and tag, and a permission boundary is attached to ensure they can't elevate permissions or access services outside EC2.

IAM Policy Evaluation Logic

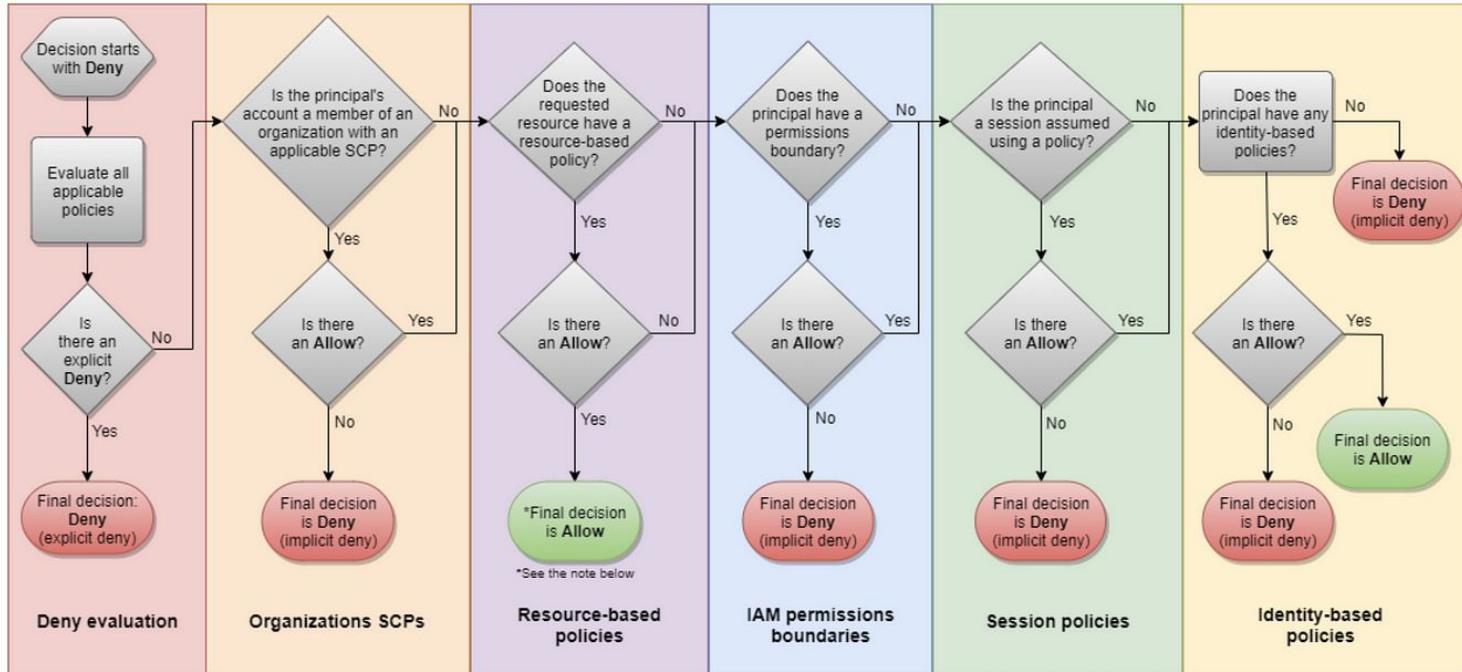
IAM **IAM Policy Evaluation Logic** is the process AWS uses to decide whether a request is allowed or denied. AWS evaluates all applicable policies (identity-based, resource-based, permissions boundaries, session policies, SCPs, and organization-level controls) in a specific order, with a key rule: explicit deny always wins.

Evaluation Logic Steps:

- By default, all requests are denied.
- An explicit allow in a policy is required to grant access.
- If any policy includes an explicit deny, it overrides all allows.
- SCPs and Permission Boundaries further limit what an identity-based policy can grant.
- This logic ensures fine-grained, layered security and prevents unauthorized actions, even if accidental permissions are granted.

Use Case:

A security team accidentally adds an Allow * policy to a junior role. However, a permission boundary and an SCP restrict that role to only read access on S3. Thanks to the evaluation logic, the actual effective permissions are limited, avoiding a security risk.



AWS IAM Identity Center

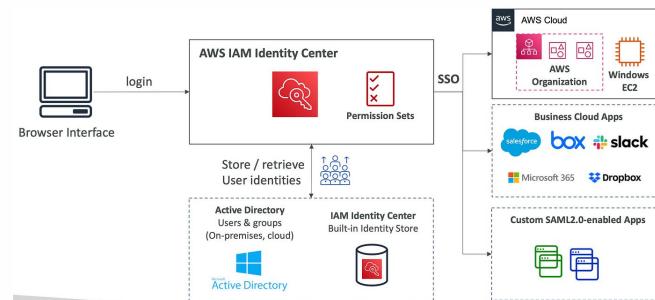
AWS **IAM Identity Center** is a centralized service for managing workforce access across AWS accounts and applications. It enables users to log in once and access multiple AWS accounts and third-party applications using Single Sign-On (SSO). It can integrate with external identity providers (IdPs) like Microsoft AD, Okta, or Azure AD using SAML 2.0 or SCIM.

Key Capabilities:

- Centralized user and group management
- Fine-grained access control using permission sets
- Integration with AWS Organizations to assign access across accounts
- Federated access using external IdPs (e.g., Azure AD)

Use Case:

A company with 50 AWS accounts wants their employees to use their existing corporate credentials to access AWS. They use IAM Identity Center with Azure AD. Each employee logs in once via the corporate portal and receives access only to the accounts and services assigned to their role, saving time and improving security.



AWS Directory Services

AWS **Directory Service** allows you to connect AWS resources to existing on-premises Microsoft Active Directory (AD), or set up a new managed directory in the cloud. It supports different directory types like:

1. **AWS Managed Microsoft AD** – fully managed AD infrastructure on AWS.
2. **AD Connector** – acts as a proxy to redirect AWS requests to your on-prem AD.
3. **Simple AD** – basic, cost-effective, standalone directory.

This service is essential for integrating AWS environments with centralized authentication systems and supporting legacy Windows-based applications.

Key Capabilities:

- Domain join for EC2 instances
- User authentication for apps like Amazon WorkSpaces or RDS for SQL Server
- Centralized access control and policy enforcement

Use Case:

An enterprise uses Microsoft AD on-prem and wants users to log into Amazon WorkSpaces using their corporate credentials. By deploying AD Connector, the AWS environment integrates with the on-prem AD, enabling seamless authentication without replicating directory data.

AWS Control Tower

AWS **Control Tower** is a service that helps you set up and govern a secure, multi-account AWS environment based on AWS best practices. It automates account provisioning, applies guardrails (preconfigured governance rules), and uses Landing Zones to establish a well-architected multi-account structure.

Control Tower integrates with:

- AWS Organizations (for account structure)
- AWS Service Catalog (for provisioning)
- CloudTrail, Config, and SNS (for logging and notifications)

Guardrails come in two types:

1. Mandatory: Always enforced (e.g., prevent deletion of log archive bucket)
2. Strongly recommended: Can be enabled/disabled by the admin

Use Case:

An enterprise wants to scale across multiple business units, each requiring separate AWS accounts with consistent security policies. Using AWS Control Tower, the cloud team can automate account creation with pre-approved settings, ensuring compliance while maintaining flexibility for each team.

Advanced Identity Exam Questions

A company is using AWS Organizations to manage multiple AWS accounts. The security team wants to ensure that no user in the “Dev” Organizational Unit (OU) can create or delete IAM roles, even if their IAM policies allow it.

What is the most effective way to enforce this restriction?

- A. Use IAM permission boundaries attached to each IAM user.
- B. Set up a service control policy (SCP) denying iam:CreateRole and iam:DeleteRole actions on the Dev OU.
- C. Configure an IAM group with a deny policy and add all users from the Dev OU.
- D. Use AWS Config rules to trigger alerts and manually roll back unauthorized changes.

A company is migrating its Windows-based applications to AWS. These applications require Active Directory (AD) for authentication. The company wants to avoid maintaining AD infrastructure on AWS but still allow EC2 instances to join the existing on-premises AD domain.

Which AWS Service option should they choose?

- A. AWS Managed Microsoft AD
- B. Simple AD
- C. AD Connector
- D. Amazon Cognito